

EP-PQC-001 — Pre-Execution Post-Quantum Cryptographic Policy Verification

ExecutionProof™ Validation Experiment — Design, Preregistration, Execution &
Decision

Remnant Fieldworks Inc. — ExecutionProof Program

October 2026

Contents

Executive summary	2
A. External-source verification report (primary sources)	3
B. Scientific-worthiness assessment	5
C. Red-team of the experiment (attack the science, not the result)	6
D. Research question (precise, falsifiable)	7
E. Threat model	8
F. Experimental policy specification	9
G. Preregistration (summary; full documents hash-sealed)	10
H. Test matrix (preregistered cases)	11
I. Evidence schema	13
J. ProofRecord requirements	14
K. Implementation architecture	15
L. Execution procedure	16
M. PASS / FAIL / INCONCLUSIVE rules (fixed before execution)	17
N. Reproducibility & package structure	18
O. Cost / time / compute	19
P. Publication / GitHub / Zenodo recommendation	20
Q. Claims we MAY make if it passes (bounded to evidence)	21
R. Claims we must NOT make (hard honesty bounds)	21
S. Follow-on experiments	22
Final decision & verdict	23

Executive summary

EP-PQC-001 asks a single, falsifiable question: can ExecutionProof physically prevent a cryptographic operation from crossing an execution boundary when the system cannot independently prove the proposed execution satisfies a defined post-quantum cryptographic (PQC) policy?

This was **not** a request to manufacture a PASS. The experiment is only worth running if the gateway (a) **independently verifies real cryptographic artifacts** rather than trusting a label, (b) **catches substitution, staleness, authority, binding, and evidence failures**, and (c) **physically blocks** the prohibited operation at the boundary.

We built a working software implementation using **real post-quantum cryptography** — ML-KEM (FIPS 203, via kyber-py 1.2.0) and ML-DSA (FIPS 204, via dilithium-py) alongside classical RSA/ECDH (via cryptography 46.0.7). The verifier does not read algorithm names and trust them; it **re-derives the algorithm family from the raw key bytes** (DER structure for classical keys; exact FIPS-203 encapsulation-key byte-lengths plus a **functional encapsulation test** for ML-KEM), **recomputes artifact hashes**, and **verifies ML-DSA-65 signatures** over evidence manifests.

The experiment was **preregistered and hash-sealed before execution**, run **blind** (a separate runner that never reads the sealed answer key), and scored afterward against the sealed key with a fixed PASS/FAIL rule.

Headline results (preserved verbatim, never massaged):

Run	Cases	Decision match	Reason match	False ALLOW	Safety-invariant violations	Verdict
EP-PQC-001 (original)	24	23/24	22/24	0	none	PARTIAL (safety-PASS, classification deviations)
EP-PQC-001R1 (remediation)	26	26/26	26/26	0	none	PASS

In both runs, the execution boundary was crossed **only** by the cases that were supposed to be allowed — never by a prohibited one. The two deviations in the original run were **both safe** (the bad operation was still blocked) and both traced to **test-design choices in our own preregistration**, not to the enforcement engine. R1 fixed the test design only — re-importing the **byte-for-byte identical engine** (impl/ep_pqc_core.py, SHA-256 506d8e0c...) — and produced a clean 26/26 PASS, which is itself evidence that the engine was correct all along.

Final decision: READY TO EXECUTE — and executed. EP-PQC-001 ran to completion; its PARTIAL result is preserved; EP-PQC-001R1 closes the two self-inflicted test-design gaps with a clean PASS.

One-sentence answer to the core question: EP-PQC-001 creates genuinely new empirical evidence for ExecutionProof — because the gateway independently re-derives the algorithm family from real cryptographic key material, recomputes hashes, and verifies real ML-DSA signatures to catch substitution, tampering, staleness, replay, authority, and provider/version failures, and the objective boundary ledger shows zero prohibited operations ever executed — rather than PQC being a mere policy label applied to capability already demonstrated elsewhere.

A. External-source verification report (primary sources)

All external PQC claims below are verified against **primary** sources, each tagged with SOURCE / REQUIREMENT / WHO IT APPLIES TO / EFFECTIVE DATE / MANDATORY-or-RECOMMENDED / CONFIDENCE. This section exists to keep the experimental policy honest: **the experiment does not claim to enforce any of these mandates.** It builds an experimental policy profile loosely inspired by them (see §F).

#	SOURCE	REQUIREMENT	APPLIES TO	EFFECTIVE DATE	STATUS	CONFIDENCE
A1	NIST FIPS 203 (ML-KEM)	ML-KEM standardized as the module-lattice KEM	Everyone standardizing on it	Finalized 13 Aug 2024	Standard (not itself a mandate)	HIGH (NIST primary)
A2	NIST FIPS 204 (ML-DSA)	ML-DSA standardized as the module-lattice signature	Everyone	Finalized 13 Aug 2024	Standard	HIGH (NIST primary)
A3	NIST FIPS 205 (SLH-DSA)	Stateless hash-based signature standardized	Everyone	Finalized 13 Aug 2024	Standard	HIGH (NIST primary)
A4	NSA CNSA 2.0 (CSA, 30 May 2025)	ML-KEM-1024 + ML-DSA-87 required; LMS/XMSS for firmware/software signing	National Security Systems (NSS)	New NSS acquisitions CNSA-2.0-compliant by 1 Jan 2027 ; exclusive use 2030-2033 by category; full NSS quantum-resistant by 2035	MANDATORY — for NSS only	HIGH-MEDIUM (NSA primary CSA; dates corroborated by multiple secondary)
A5	NSA CNSA 2.0	Excludes SLH-DSA and QKD for NSS use	NSS	—	MANDATORY exclusion for NSS	MEDIUM (NSA FAQ/CSA)
A6	CNSSP-15 (2025 revision)	Incorporates CNSA 2.0 as the policy vehicle for NSS	NSS owners	2025	MANDATORY for NSS	MEDIUM
A7	NSM-10 (2022)	Mitigate quantum risk to NSS by 2035	US Government NSS	2035 target	MANDATORY (policy memorandum)	HIGH

#	SOURCE	REQUIREMENT	APPLIES TO	EFFECTIVE DATE	STATUS	CONFIDENCE
A8	CISA / NSA / NIST joint guidance	Maintain a cryptographic inventory ; plan migration; prioritize HNDL-exposed data	Critical infrastructure & commercial (advisory)	Ongoing	RECOMMENDED (not statutory for commercial)	DECH
A9	Industry (Google, Apple, Cloudflare, Microsoft)	Hybrid KEM (e.g., X25519+ML-KEM) deployment in TLS/messaging	Their own products	2023–2026 rollouts	Vendor practice, not a mandate	HIGH (vendor announcements)

Load-bearing cautions (explicitly honored in the design):

- **Do not translate an NSS requirement into a universal commercial mandate.** CNSA 2.0 / CNSSP-15 bind National Security Systems. They are **not** a statutory obligation on general commercial systems. Our policy profile is labeled experimental, not “CNSA 2.0 compliant.”
- **Do not translate a recommendation into a statutory deadline.** Cryptographic inventory and migration planning are **recommended** best practice for commercial entities, not law.
- **“Harvest now, decrypt later” (HNDL)** is the real driver: long-lived confidential data is exposed today even though cryptographically-relevant quantum computers do not yet exist publicly. This justifies pre-execution prevention rather than post-hoc detection.
- The 2027/2030/2035 dates originate in the **NSA CSA of 30 May 2025** (and NSM-10), not the 2022 CNSA 2.0 FAQ. We attribute them to the correct document.

B. Scientific-worthiness assessment

Claim under scrutiny: is EP-PQC-001 genuinely new evidence, or is “PQC” just another label bolted onto capability already shown by prior ExecutionProof work (ARK boundary series, WITNESS quantum-nonce, QT verification)?

Why it is worth running (and what would have made it worthless):

Dimension	Worthless version (rejected)	Worthwhile version (built)
Artifact handling	Trust <code>claimed_algorithm == "ML-KEM"</code> string	Independently derive family from key bytes; functional ML-KEM <code>encaps()</code> test
Evidence	Accept an unsigned JSON “attestation”	Verify real ML-DSA-65 signature over a manifest; recompute artifact SHA-256
Policy	Hard-code if algo in <code>APPROVED: allow</code>	Multi-stage: authority → evidence → constraint, each independently falsifiable
Enforcement	Log “would block”	Objective boundary ledger: protected ML-KEM encapsulation physically runs only on ALLOW
Failure realism	Only happy-path cases	24-26 cases incl. substitution, tamper, stale, replay, downgrade, untrusted provider, custody

Distinct from prior series: ARK tested authorization-boundary representation on hardware; WITNESS/BELLWETHER/CHRONO/TRINITY bind quantum witnesses into a nonce; QT tests tail-aware verification. **None of them re-derives a cryptographic algorithm family from raw key material and gates a real PQC operation on that independent derivation.** That specific capability — “the label on evidence is not the provenance of evidence,” applied to cryptographic artifacts — is what EP-PQC-001 newly demonstrates.

Worthiness verdict: WORTH RUNNING. The experiment is non-trivial precisely because the verifier can fail — and in the original run it exposed two gaps in our own test design, which is exactly what an honest preregistered experiment is for.

C. Red-team of the experiment (attack the science, not the result)

We deliberately tried to break the scientific validity of EP-PQC-001. Each objection and its disposition:

1. **“You’re just matching strings in disguise.”** — Refuted. PQC-04/PQC-16 substitute the algorithm while keeping the claim fixed; the verifier catches them by **byte-level derivation** and functional testing, not by comparing claim strings. Removing the derivation code makes these cases fail — proving the derivation is load-bearing.
2. **“The authority stage short-circuits everything, so evidence is untested.”** — This is the real PQC-04 deviation. Honestly disclosed. R1 isolates the evidence path by making the requester the true custodian; the evidence stage then independently fires. The engine was never wrong; our case was mis-targeted.
3. **“Blinding is fake — same author wrote generator and verifier.”** — Conceded. Blinding is **procedural**, not third-party. We mitigate by (a) a separate sealed answer key the runner never reads, (b) hash-sealing the preregistration before execution, and (c) an objective boundary ledger that cannot be argued with. We state this limitation plainly rather than overclaim.
4. **“A software gateway isn’t a real boundary.”** — Conceded scope. This is a **simulation**: no HSM/TPM, no live TLS, no kernel enforcement. The ledger proves logical prevention, not physical hardware prevention. EP-PQC-005/006 (see §S) are the hardware/TLS follow-ons.
5. **“You chose policies that guarantee a PASS.”** — Refuted by construction: 20 of 26 R1 cases are expected non-ALLOW, and the original run did **not** fully PASS. The policies are adversarial to the happy path.
6. **“Reason codes are cosmetic.”** — Partly conceded: the safety verdict depends only on decision + boundary, not reason text. Reason match is a **stricter** secondary bar; the two original deviations were reason-adjacent, not decision errors.

Red-team conclusion: the experiment’s safety claim (zero prohibited crossings) is robust; the classification claim is honestly bounded by procedural (not independent) blinding and by simulation scope.

D. Research question (precise, falsifiable)

RQ: Given a proposed cryptographic operation accompanied by (i) an algorithm/parameter claim, (ii) a real public-key artifact, and (iii) a signed evidence manifest, will an ExecutionProof gateway **physically withhold execution** (HOLD or DENY, with no boundary crossing) in every case where it cannot independently verify that the artifact, its provenance, and the requester’s authority jointly satisfy a defined PQC policy — while allowing execution **only** when all three independently verify?

Falsification condition: a single **False ALLOW** — any case expected to be HOLD/DENY whose protected operation nonetheless crosses the boundary — falsifies the safety claim.

E. Threat model

Adversary capabilities assumed:

- Can submit execution requests with **arbitrary algorithm/parameter claims** (substitution).
- Can supply **mismatched or tampered artifacts** (wrong key, flipped bytes).
- Can **replay** previously valid evidence (cross-request and same-request nonce replay).
- Can present **stale** evidence, **untrusted-provider** evidence, or **below-minimum versions**.
- Can attempt operations **without authority** or **without custody** of the referenced key.
- Can attempt **parameter downgrade** (ML-KEM-512 where $\geq 768/1024$ is required) and **hybrid smuggling** where a PQC-only policy forbids it.

Adversary limitations assumed:

- Cannot forge an **ML-DSA-65 signature** without the authority's secret key (standard EUF-CMA assumption; assumed, not proven here).
- Cannot find a **SHA-256 collision** for an artifact.
- Cannot modify the **enforcement engine, policy files, or sealed answer key** (integrity by hash).

Out of scope (honestly declared): side-channels, HSM/TPM compromise, supply-chain compromise of the crypto libraries themselves, and the mathematical security of ML-KEM/ML-DSA (we use them; we do not prove them).

F. Experimental policy specification

Two **experimental** policy profiles (JSON, hash-committed). They are **not** “CNSA 2.0 compliant” labels — they are deliberately-constructed test policies inspired by the migration landscape.

EP-PQC-POLICY-001-A (a1642a74...) — transitional:

- Allowed KEM family: ML-KEM, rank ≥ 3 (i.e., ML-KEM-768 and ML-KEM-1024).
- Hybrid (classical+PQC) **permitted**.
- Trusted providers: {kyber-py}; minimum version 1.2.0.
- Legacy classical algorithms $\Rightarrow$ CONSTRAINT_LEGACY_ALGORITHM (HOLD).

EP-PQC-POLICY-001-B (99663b47...) — rollover / strict:

- Allowed: ML-KEM-1024 **only**; **PQC-only** (hybrid forbidden $\Rightarrow$ CONSTRAINT_HYBRID_NOT_PERMITTED).
- Everything below 1024 $\Rightarrow$ CONSTRAINT_ALGORITHM_NOT_ALLOWED / CONSTRAINT_PARAM_DOWNGRADE.

Policies define **constraint** rules only. They **cannot** override the structural invariants (authority/evidence precede constraint), and they contain **no** “approved algorithm” shortcut that bypasses independent artifact derivation.

G. Preregistration (summary; full documents hash-sealed)

The full preregistration was written and **hash-sealed before execution**:

Document	SHA-256
EP-PQC-001_PREREGISTRATION.md	931508ea1de6f87623613ca0a247752da666769671dc6fbd5ca50
R1/EP-PQC-001R1_PREREGISTRATION.md	fce80a02f21837e113f24942395ada532f4023b818354d987cfb8

Each preregistration commits: the invariants, the PASS/FAIL rule, the blinding procedure, the case count, and the **frozen hashes** of challenges, sealed answer key, authority public key, and the enforcement engine. The engine hash is **identical** across both runs — the committed proof that R1 changed the test, not the enforcement logic.

H. Test matrix (preregistered cases)

26 R1 cases (24 in the original run; R1 adds the PQC-17a/b replay pair and re-targets PQC-04/12). Decision + reason are the sealed expectations; “crossed” is the objective boundary outcome.

Case	Policy	Dimension tested	Expected decision	Expected primary reason	Crossed?
PQC-01	A	Fully valid PQC	ALLOW	ALL_STAGES_PASS	yes
PQC-02	A	Honest legacy RSA under PQC policy	HOLD	CONSTRAINT_LEGACY_ALGORITHM	no
PQC-03	A	No evidence	HOLD	EVID_MISSING	no
PQC-04	A	Alg claim mismatch (RSA as ML-KEM), custodian correct	HOLD	EVID_ALGORITHM_CLAIM_MISMATCH	no
PQC-05	A	Unauthorized requester	DENY	AUTH_REQUESTER_NOT_AUTHORIZED	no
PQC-06	A	Tampered key bytes	HOLD	EVID_HASH_MISMATCH	no
PQC-07	A	Stale evidence	HOLD	EVID_STALE	no
PQC-08	A	Evidence bound to a different key	HOLD	EVID_ARTIFACT_BINDING_MISMATCH	no
PQC-09	A	Approved hybrid under A	ALLOW	ALL_STAGES_PASS	yes
PQC-10	B	Hybrid under PQC-only B	HOLD	CONSTRAINT_HYBRID_NOT_PERMITTED	no
PQC-11a	A	Rollover part 1 (allowed under A)	ALLOW	ALL_STAGES_PASS	yes
PQC-11b	B	Rollover part 2 (forbidden under B)	HOLD	CONSTRAINT_ALGORITHM_NOT_ALLOWED	no
PQC-12	A	Cross-request replay	HOLD	EVID_REQUEST_BINDING_MISMATCH	no
PQC-13	A	Untrusted provider	HOLD	CONSTRAINT_UNTRUSTED_PROVIDER	no
PQC-14	A	Parameter downgrade (512)	HOLD	CONSTRAINT_PARAM_DOWNGRADE	no
PQC-15	A	Version below minimum	HOLD	CONSTRAINT_VERSION_BELOW_MINIMUM	no
PQC-16	A	Substitution 1024-claim vs 768-artifact	HOLD	EVID_ALGORITHM_CLAIM_MISMATCH	no
PQC-17a	A	Valid request; consumes nonce	ALLOW	ALL_STAGES_PASS	yes

Case	Policy	Dimension tested	Expected decision	Expected primary reason	Crossed?
PQC-17b	A	Same-request nonce replay	HOLD	EVID_NONCE_REPLAY	
NC-00	A	Baseline (must ALLOW)	ALLOW	ALL_STAGES_PASSED	
NC-01	A	Only authority changed	DENY	AUTH_REQUESTER_NOT_AUTHORIZED	
NC-02	A	Only provider changed	HOLD	CONSTRAINT_UNTRUSTED_PROVIDER	
NC-03	B	Only policy version changed	HOLD	CONSTRAINT_ALGORITHM_NOT_ALLOWED	
NC-04	A	Only freshness changed	HOLD	EVID_STALE	no
NC-05	A	Only algorithm changed	HOLD	CONSTRAINT_PARAM_DOWNGRADE	
NC-06	A	Only artifact binding changed	HOLD	EVID_ARTIFACT_BINDING_MISMATCH	

Negative controls (NC-01...NC-06) each change exactly **one** dimension from the NC-00 baseline, isolating which stage flips. All six behaved as designed.

I. Evidence schema

Every request carries an **evidence bundle**:

```
bundle = {  
  "manifest": {  
    "artifact_ref", "artifact_sha256", "issued_at", "nonce",  
    "request_id", "provider", "version"  
  },  
  "manifest_sig": <ML-DSA-65 signature over canonical manifest bytes>,  
  "public_key_b64": <the actual public-key artifact being attested>  
}
```


Independent checks performed by the verifier (none trust a label):

1. manifest_sig verifies under the authority's ML-DSA-65 public key → else EVID_SIGNATURE_INVALID.
2. sha256(public_key_b64) == manifest.artifact_sha256 → else EVID_HASH_MISMATCH (tamper).
3. manifest.request_id == request.request_id → else EVID_REQUEST_BINDING_MISMATCH (replay).
4. manifest.nonce == request.nonce and nonce unused → else EVID_NONCE_* (binding/replay).
5. manifest.artifact_ref == request.artifact_ref → else EVID_ARTIFACT_BINDING_MISMATCH.
6. now - issued_at ≤ freshness_window → else EVID_STALE.
7. **Independent algorithm derivation** of public_key_b64 must match claimed_algorithm → else EVID_ALGORITHM_CLAIM_MISMATCH.

J. ProofRecord requirements

Every decision — ALLOW, HOLD, or DENY — emits a **ProofRecord** JSON (26 files in R1/proofrecords/). Each record contains: the request, the per-stage results (authority / evidence / constraint) with their reason codes, the independently-**derived** algorithm, the policy version and its source hash, the gateway status, the boolean **boundary_crossed**, the environment fingerprint, and a self-record_hash. ProofRecords are the auditable, tamper-evident trail: a reviewer can replay any record against the engine and reproduce the identical decision and hash.

K. Implementation architecture



- `inspect_artifact()` derives the family from raw bytes: classical keys via cryptography DER parse (RSA/EC/Ed25519); ML-KEM via **exact FIPS-203 encapsulation-key byte-lengths** (800/1184/1568 for 512/768/1024) **plus a functional `encaps()` call** that must succeed.
- ExecutionGateway maintains the **objective boundary ledger**: the protected operation (a real ML-KEM encapsulation) executes **only** inside `cross()`, which is reachable **only** on ALLOW.
- Core engine: `impl/ep_pqc_core.py`, 496 lines, SHA-256 506d8e0c... — **identical** across both runs.

L. Execution procedure

1. `generator/generate_cases.py` (and `R1/generator/generate_cases_r1.py`) build real artifacts, sign manifests, and write `challenges.json` (no labels) + `answer_key.sealed.json` (sealed).
2. The preregistration is written and its SHA-256 recorded **before** any run.
3. `run_experiment.py` (and `R1/run_experiment_r1.py`) run **blind** — reading only challenges, policy, and the authority public key — emitting `raw_decisions.json` + `ProofRecords`.
4. `score.py` (and `R1/score_r1.py`) compare against the sealed key **after** the run and apply the fixed PASS/FAIL rule.
5. A deterministic evaluation clock ($T_0 = 1_760_000_000$) makes freshness reproducible.

M. PASS / FAIL / INCONCLUSIVE rules (fixed before execution)

- **SAFETY PASS** $\iff$ `False_ALLOW == 0` **and** no non-ALLOW expectation crosses the boundary.
- **EXACT-CLASSIFICATION PASS** $\iff$ every case matches on **both** decision and primary reason code.
- **VERDICT = PASS** $\iff$ SAFETY PASS **and** EXACT-CLASSIFICATION PASS.
- **VERDICT = PARTIAL** $\iff$ SAFETY PASS but not EXACT-CLASSIFICATION PASS (deviations preserved).
- **VERDICT = FAIL** $\iff$ not SAFETY PASS.
- **INCONCLUSIVE** is reserved for runs that cannot be scored (environment/crypto-library failure), **not** for safe classification deviations. A FAIL is **never** reclassified as INCONCLUSIVE.

The worst error class is **False ALLOW** (a prohibited operation executing). Both runs: **0**.

N. Reproducibility & package structure

```
RF_EP_PQC_001/
├── impl/ep_pqc_core.py          # enforcement engine (frozen, hashed)
├── policy/EP-PQC-POLICY-001-A.json # transitional policy (hashed)
├── policy/EP-PQC-POLICY-001-B.json # rollover policy (hashed)
├── generator/generate_cases.py   # original generator
│   ├── challenges.json          # 24 blind challenges
│   └── answer_key.sealed.json   # sealed expectations
├── evidence/authority_pubkey.json # authority ML-DSA-65 public key
├── run_experiment.py / score.py  # blind runner + scorer
├── EP-PQC-001_PREREGISTRATION.md # hash-sealed prereg (+ .txt)
├── results/ (raw_decisions.json, scored_results.json)
├── proofrecords/ (24 hashed records)
└── R1/                          # remediation: same engine, corrected test design
    ├── generator/generate_cases_r1.py (+ challenges.json, answer_key.sealed.json)
    ├── evidence/authority_pubkey.json
    ├── run_experiment_r1.py / score_r1.py
    ├── EP-PQC-001R1_PREREGISTRATION.md (+ .txt)
    ├── results/ (raw_decisions.json, scored_results.json)
    └── proofrecords/ (26 hashed records)
```

Environment: Python 3.11.6; cryptography 46.0.7; kyber-py 1.2.0 (ML-KEM/FIPS 203); dilithium-py (ML-DSA/FIPS 204). Deterministic clock. To reproduce: re-run the runner + scorer against the frozen challenges; decisions and ProofRecord hashes are deterministic. (Re-running a generator mints fresh authority keys and therefore new challenge/answer-key hashes — the committed artifacts are **generate-once-and-freeze.**)

O. Cost / time / compute

Negligible. The entire experiment runs in **seconds** on a single CPU core with no GPU, no network, and no paid services. Real ML-KEM/ML-DSA operations are millisecond-scale. Reproduction cost $\approx$ \$0. This is a **software simulation**; a hardware/HSM follow-on (\$\$) would carry real device cost.

P. Publication / GitHub / Zenodo recommendation

Recommendation: HOLD for Derek's explicit written GO before any public release or DOI minting (consistent with RF policy for QT-TR and prior series). When/if released, the package is self-contained and reproducible. Suggested framing: "a preregistered software validation of pre-execution PQC policy enforcement" — never "CNSA 2.0 compliance," never "proven secure." No DOI should be minted and no preprint posted without Derek's written authorization.

Q. Claims we MAY make if it passes (bounded to evidence)

- ExecutionProof can **independently derive** a cryptographic algorithm family from raw key material and **gate a real PQC operation** on that derivation — demonstrated on real ML-KEM/ML-DSA/RSA/ECDH artifacts.
- In a preregistered, blind, hash-sealed software experiment, the gateway **physically withheld** execution for **every** case it could not independently verify (26/26 in R1), with **0 False ALLOW** and **0 boundary violations**.
- The gateway catches **substitution, tampering, staleness, cross-request replay, nonce replay, authority/custody failure, provider and version violations, parameter downgrade, and hybrid smuggling** — each via an independent check, not a label comparison.

R. Claims we must NOT make (hard honesty bounds)

- **Not** “CNSA 2.0 / FIPS compliant” — we build an experimental policy profile, not a certified compliance control.
- **Not** “proves ML-KEM/ML-DSA secure” — we use these primitives; we do not establish their cryptographic security.
- **Not** “hardware / HSM / TLS / production enforcement” — this is a **software simulation** with no HSM/TPM, no live handshake, no kernel/hardware boundary.
- **Not** “independently validated” — blinding is **procedural** (single author), not third-party.
- **Not** “first ever” / “patented” / “unbreakable” — no such claim is supported by this evidence.
- The experiment demonstrates **successful policy enforcement in simulation**, which is distinct from **successful cryptographic migration** or **successful production execution enforcement**.

S. Follow-on experiments

ID	Title	What it would add beyond EP-PQC-001
EP-PQC-002	Independent third-party blinding	A separate party generates the sealed set; removes the procedural-blinding caveat.
EP-PQC-003	Adversarial fuzzing of artifacts	Automated mutation of key bytes/manifests to probe for any False ALLOW at scale.
EP-PQC-004	Full CNSA-2.0-shaped NSS profile	A policy profile that does mirror CNSA 2.0 (ML-KEM-1024 + ML-DSA-87 + LMS/XMSS), still labeled experimental.
EP-PQC-005	HSM / TPM-backed gateway	Move the boundary into real hardware; test physical (not logical) prevention.
EP-PQC-006	Live TLS handshake enforcement	Gate a real hybrid TLS 1.3 key exchange; prevent a non-compliant handshake from completing.
EP-PQC-007	Crypto-agility rollover under load	A/B policy rollover (transitional→strict) during continuous traffic; measure zero-gap enforcement.

Final decision & verdict

DECISION: READY TO EXECUTE — and executed.

EP-PQC-001 was designed, preregistered, hash-sealed, run blind, and scored. Its **official verdict is PARTIAL (safety-PASS, classification deviations)** and is preserved verbatim. The remediation run **EP-PQC-001R1** — using the **byte-for-byte identical engine** and correcting only two self-inflicted test-design choices plus adding an explicit nonce-replay case — produced a clean **PASS (26/26 decision, 26/26 reason, 0 False ALLOW, 0 boundary violations)**.

One-sentence answer to the commissioned question:

EP-PQC-001 creates genuinely new empirical evidence for ExecutionProof — because the gateway independently re-derives the algorithm family from real cryptographic key material, recomputes hashes, and verifies real ML-DSA-65 signatures to catch substitution, tampering, staleness, replay, authority, custody, provider, version, downgrade, and hybrid-smuggling failures, and the objective boundary ledger records zero prohibited operations ever executed — so PQC here is a physically-enforced, independently-verified capability, not merely a policy label applied to capability already demonstrated elsewhere.

Honesty footer. Software-only simulation (no HSM/TPM/live TLS). Blinding is procedural, not third-party. This experiment uses, but does not prove the security of, ML-KEM/ML-DSA. The experimental policy profiles are inspired by — but are **not** — CNSA 2.0 / FIPS compliance controls. All failures are preserved; no goalpost was moved; the original PARTIAL result was never overwritten.